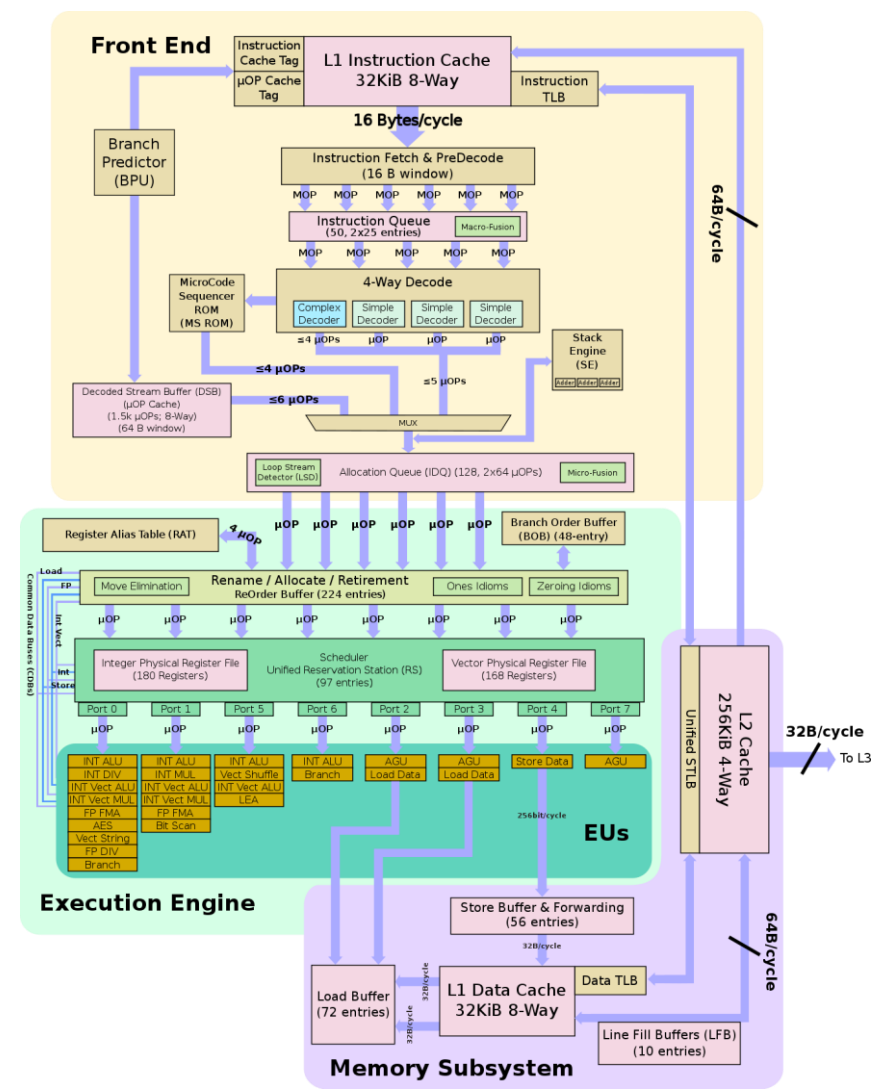


谢天


单核篇

# CPU Performance



# Disclaimer

## 声明



1. 分享内容基于个人的经验，测试数据都是特定情况下的数据，不具有通用性
2. Opinions are my own, 我的观点不代表FF引擎组 or 公司的观点
3. 时间有限，只能分享一些我觉得有趣的方面，欢迎提问交流/会后交流

# Performance Matters



**Sebastian Aaltonen**  
@SebAaltonen

...

Example of modern user-generated content in HypeHype: This scene has thousands of skinned characters, thousands of physics objects and over 10,000 audio sources. It runs at 60 fps on 99\$ phone. It doesn't look fancy, but it would choke Unity even on high end PC.



**Sebastian Aaltonen** @SebAaltonen · Nov 29

...

Worth noting: Above example has no tricks at all. Each character's bone animation and skinning runs at 60Hz. There's no animation sharing. The trick is to optimize the code. Then you don't need to limit character count or run animation at 10 fps.


12

3

226

9.3K

<https://x.com/SebAaltonen/status/1862397887541387424>



## Performance Tuning

### Low-level optimizations (On Single Core)


takes into account the details of the underlying hardware capabilities.

### High-level optimizations

more about application-level logic, algorithms, and data structures.

## The free lunch is over.

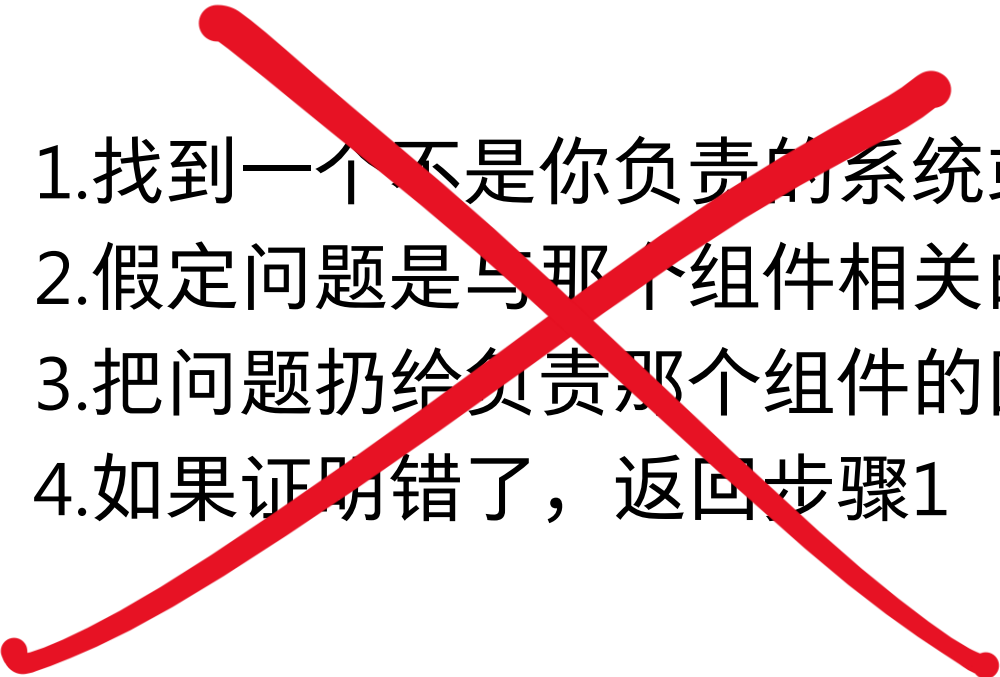
50 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2021 by K. Rupp

## 性能分析方法

### 责怪他人法

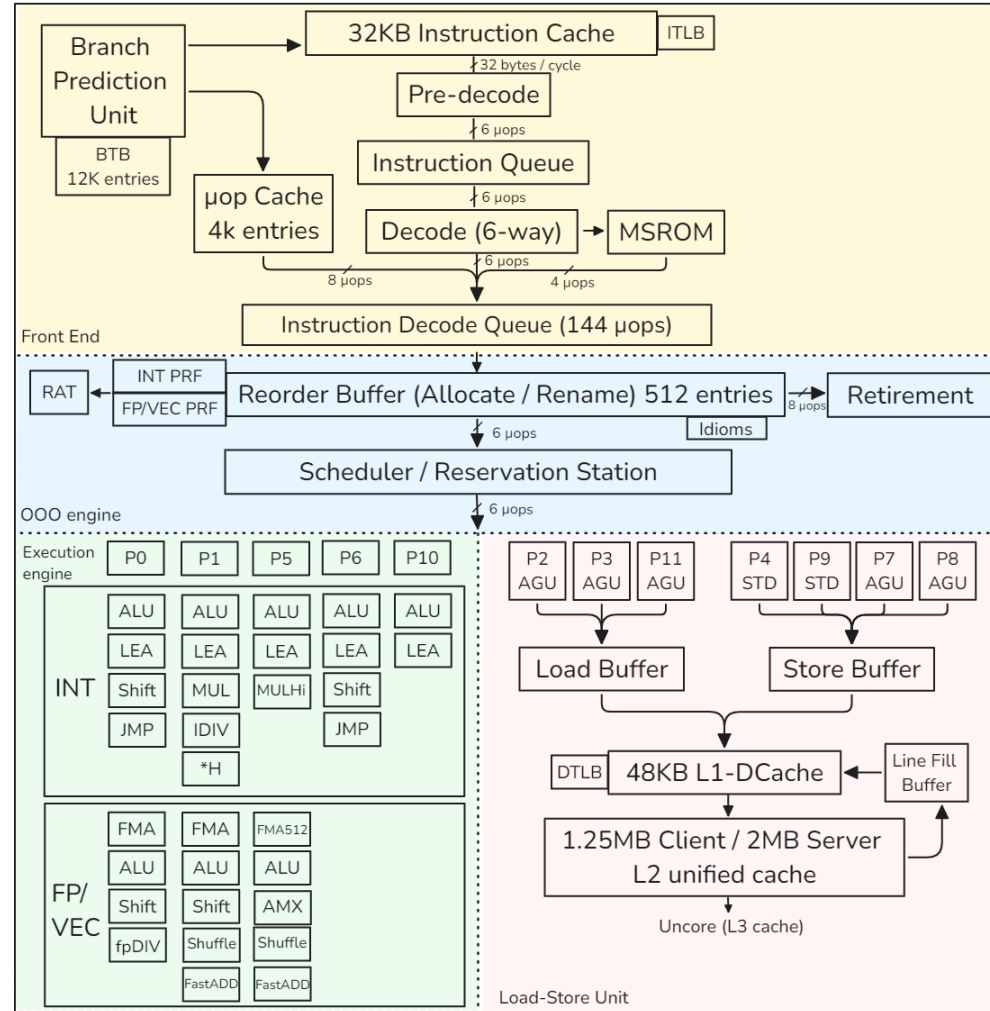
- 
1. 找到一个不是你负责的系统或环境的组件
  2. 假定问题是与那个组件相关的
  3. 把问题扔给负责那个组件的团队
  4. 如果证明错了，返回步骤1

# 性能分析方法

## Identifying CPU bottlenecks

# Top-down Microarchitecture Analysis Method

# Microarchitecture: CPU simplified view

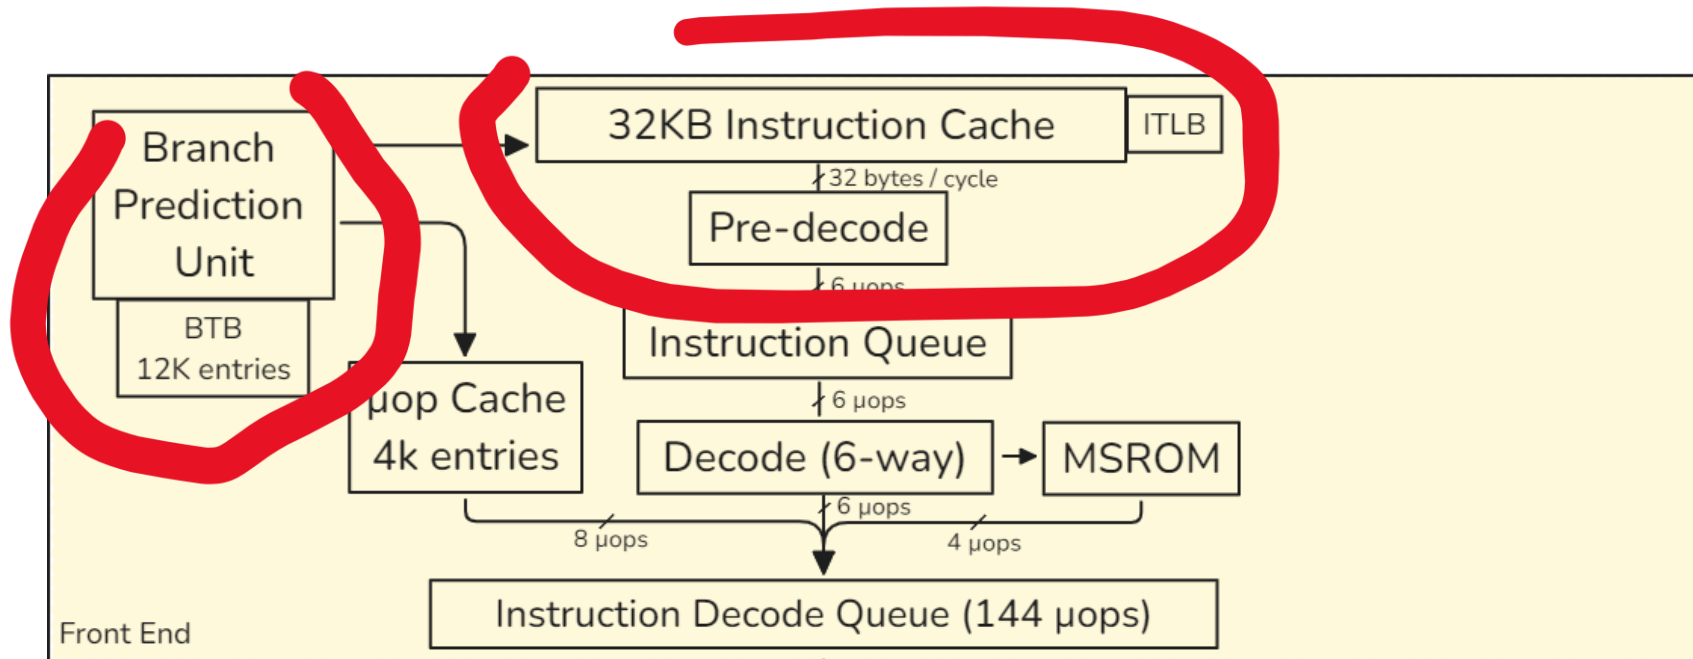


the implementation of Intel 12th-generation core, Golden Cove



## Front End

### Fetch and decode instructions from memory



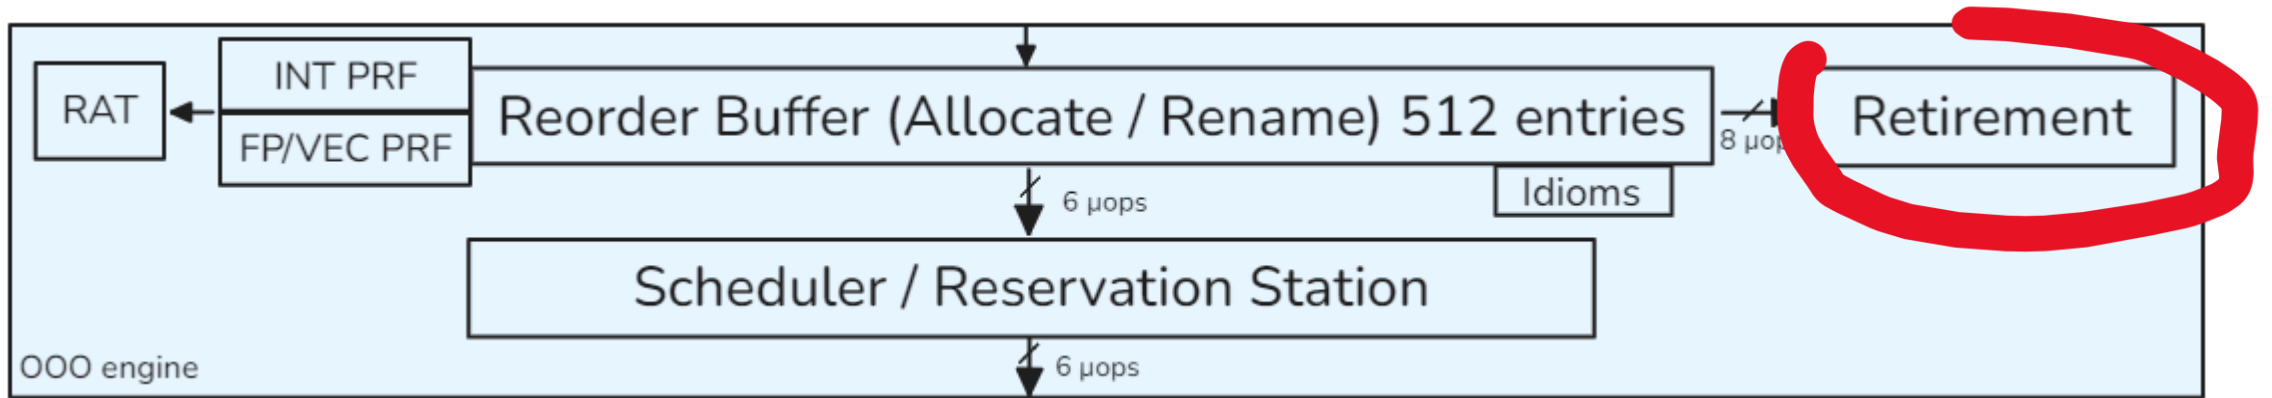
x86: translate CISC instructions into simple  $\mu$ ops => TMAM based on  $\mu$ ops/uop

The CPU Frontend **fetches** 32 bytes per cycle of x86 instructions from the L1 I-cache => **Potential stall!**

The **BPU** predicts the target of all branch instructions based on prediction => **Potential stall!**

## Back End

an OOO engine that executes instructions and stores results

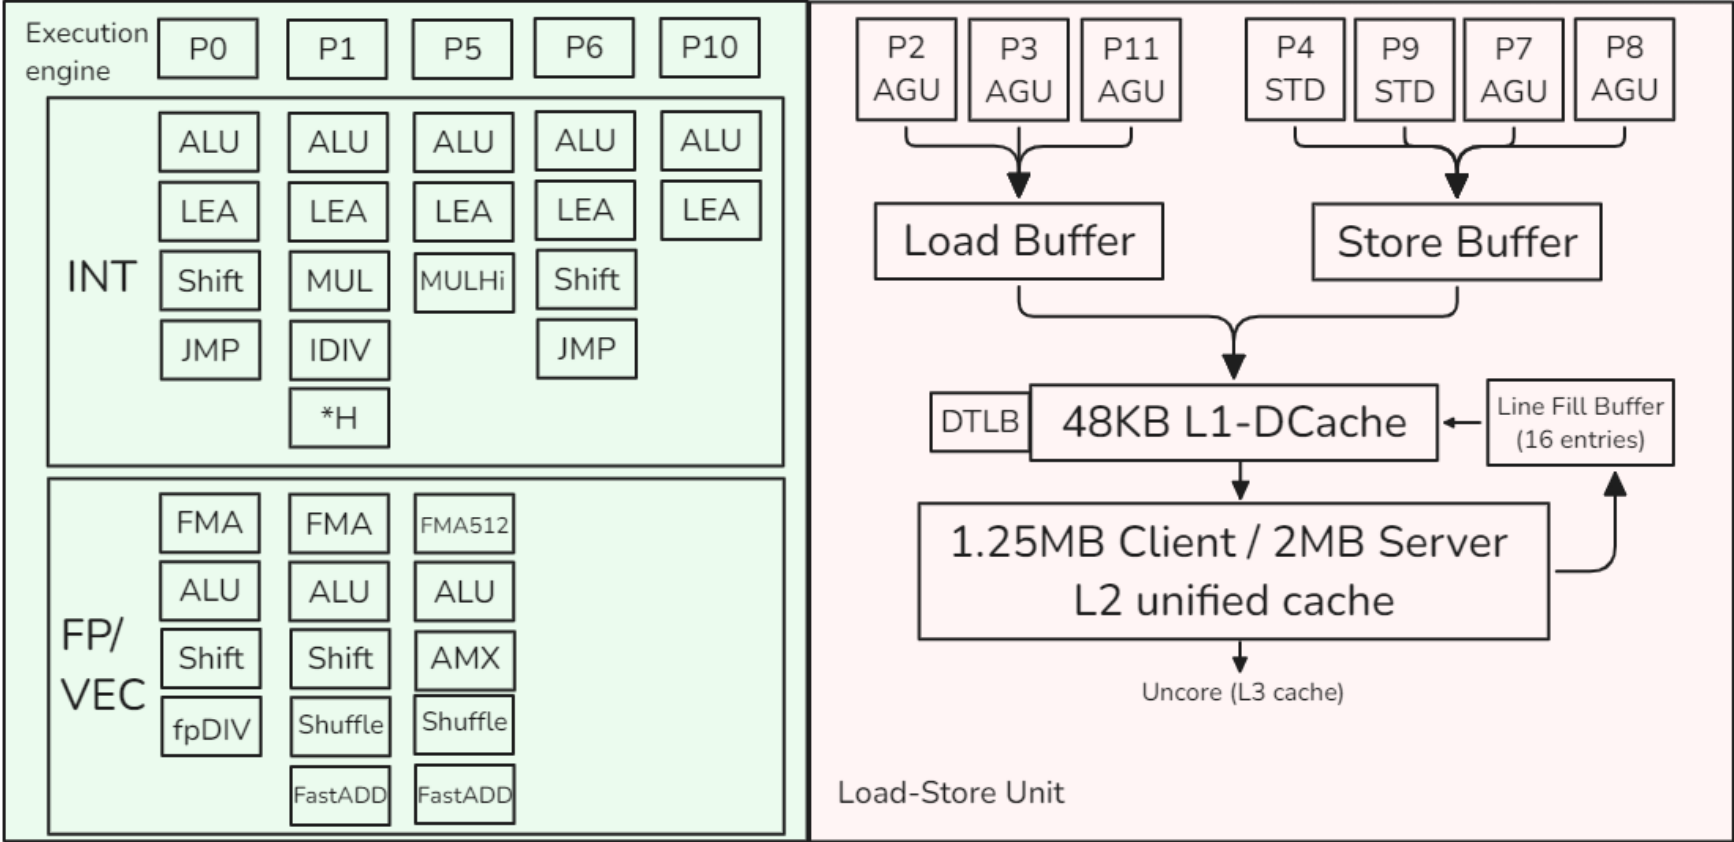


ReOrder Buffer: Register renaming, allocates execution resources, tracks speculative execution.

The Scheduler / Reservation Station: tracks the availability of all resources for a given  $\mu$ op and dispatches the  $\mu$ op to an execution port once it is ready.

# Back End

## The execution engine and the Load-Store unit



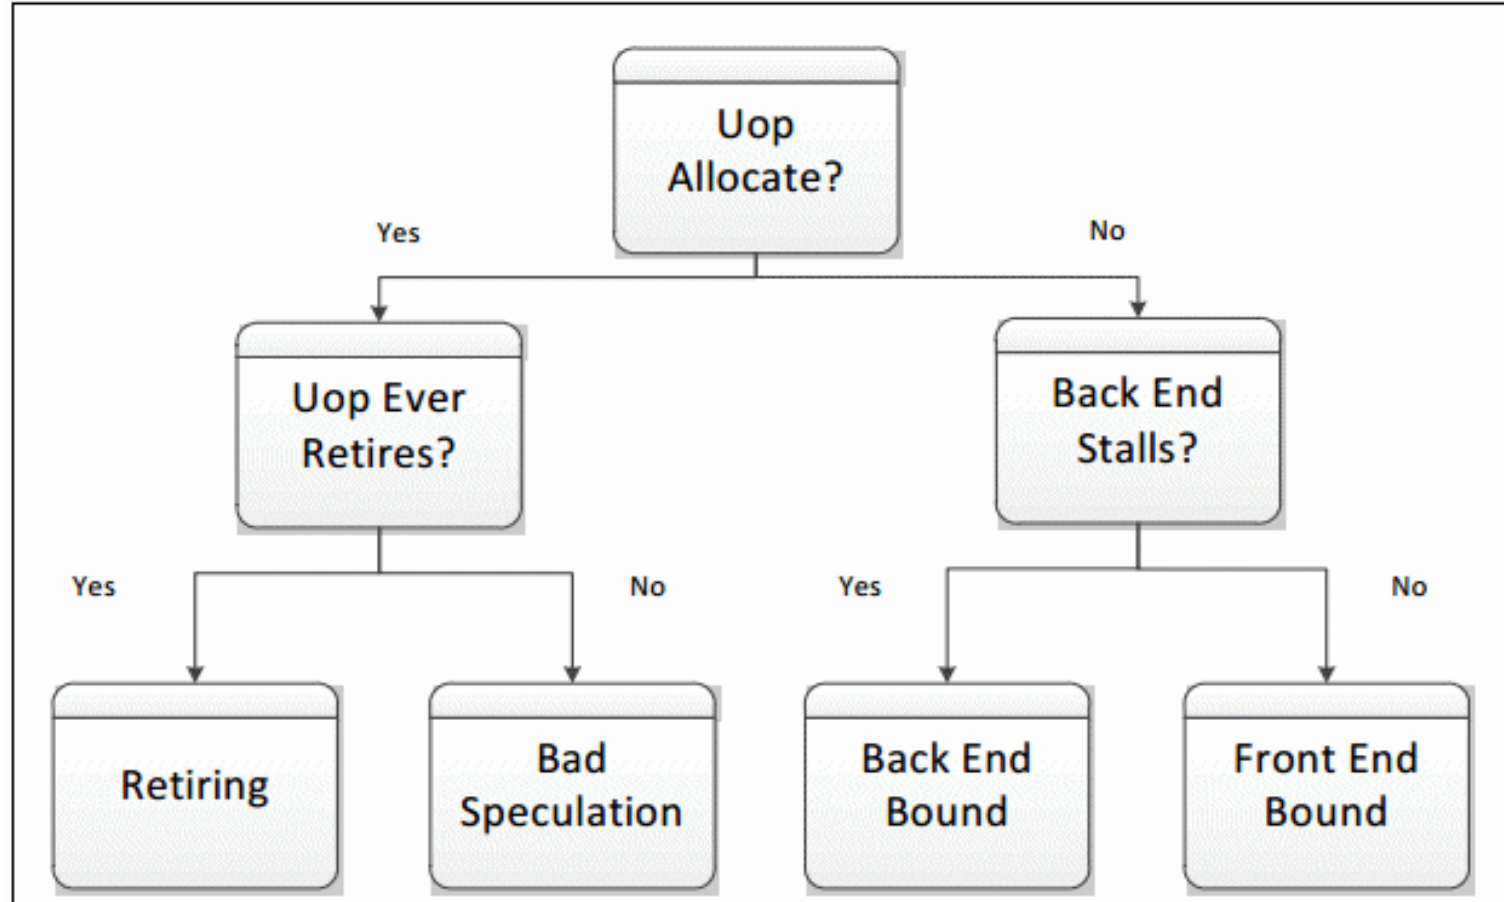
### 12 execution ports

When a scheduler has to dispatch two operations that require the same execution port, one of them will have to be delayed.

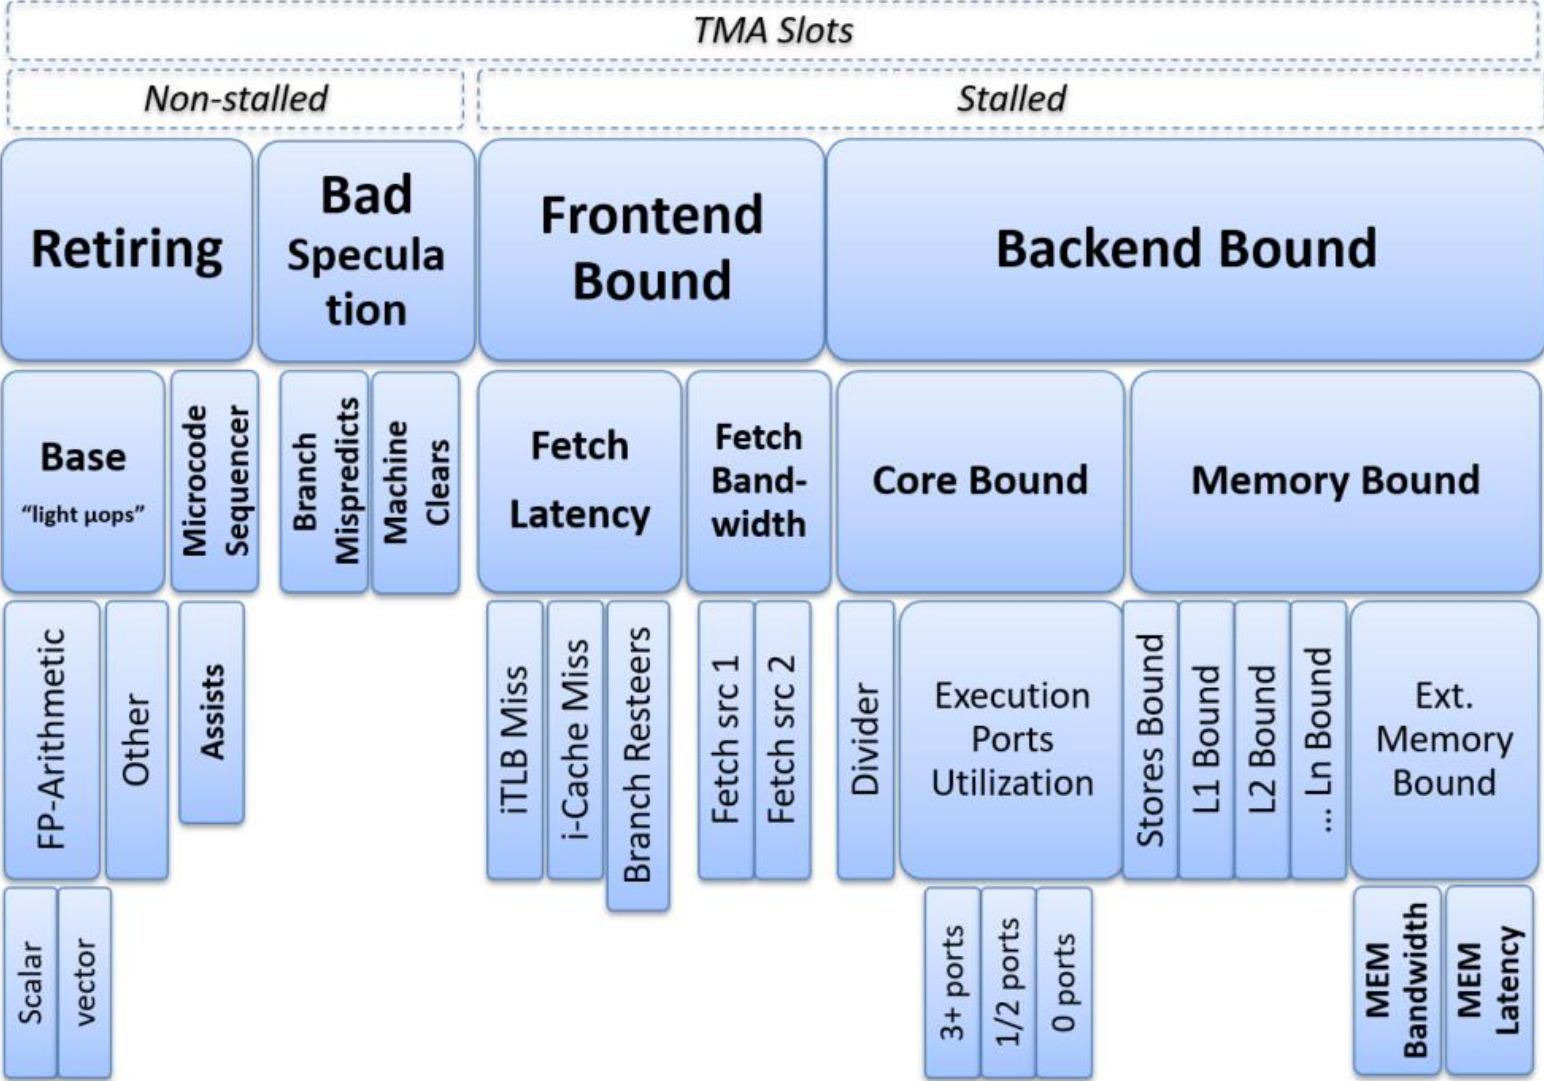
Potential stall!

Potential stall!

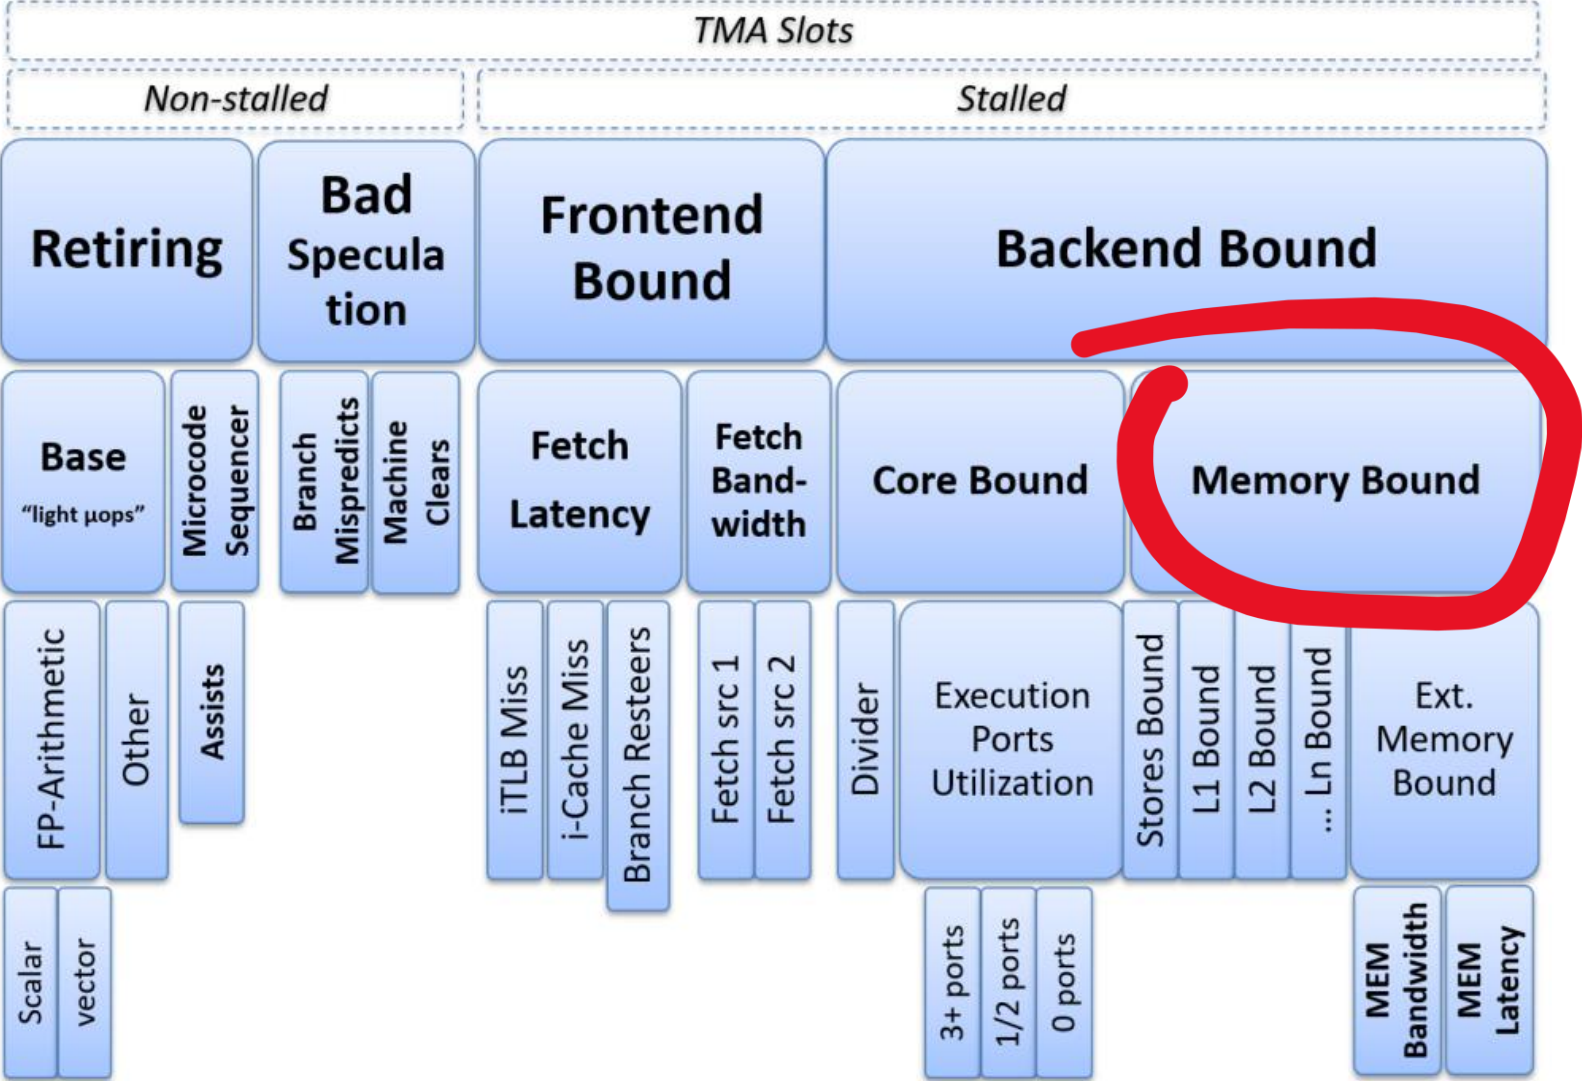
## TMAM: Pipeline slot




# TMA hierarchy of performance bottlenecks



# Four Corners



# Processor Memory Gap

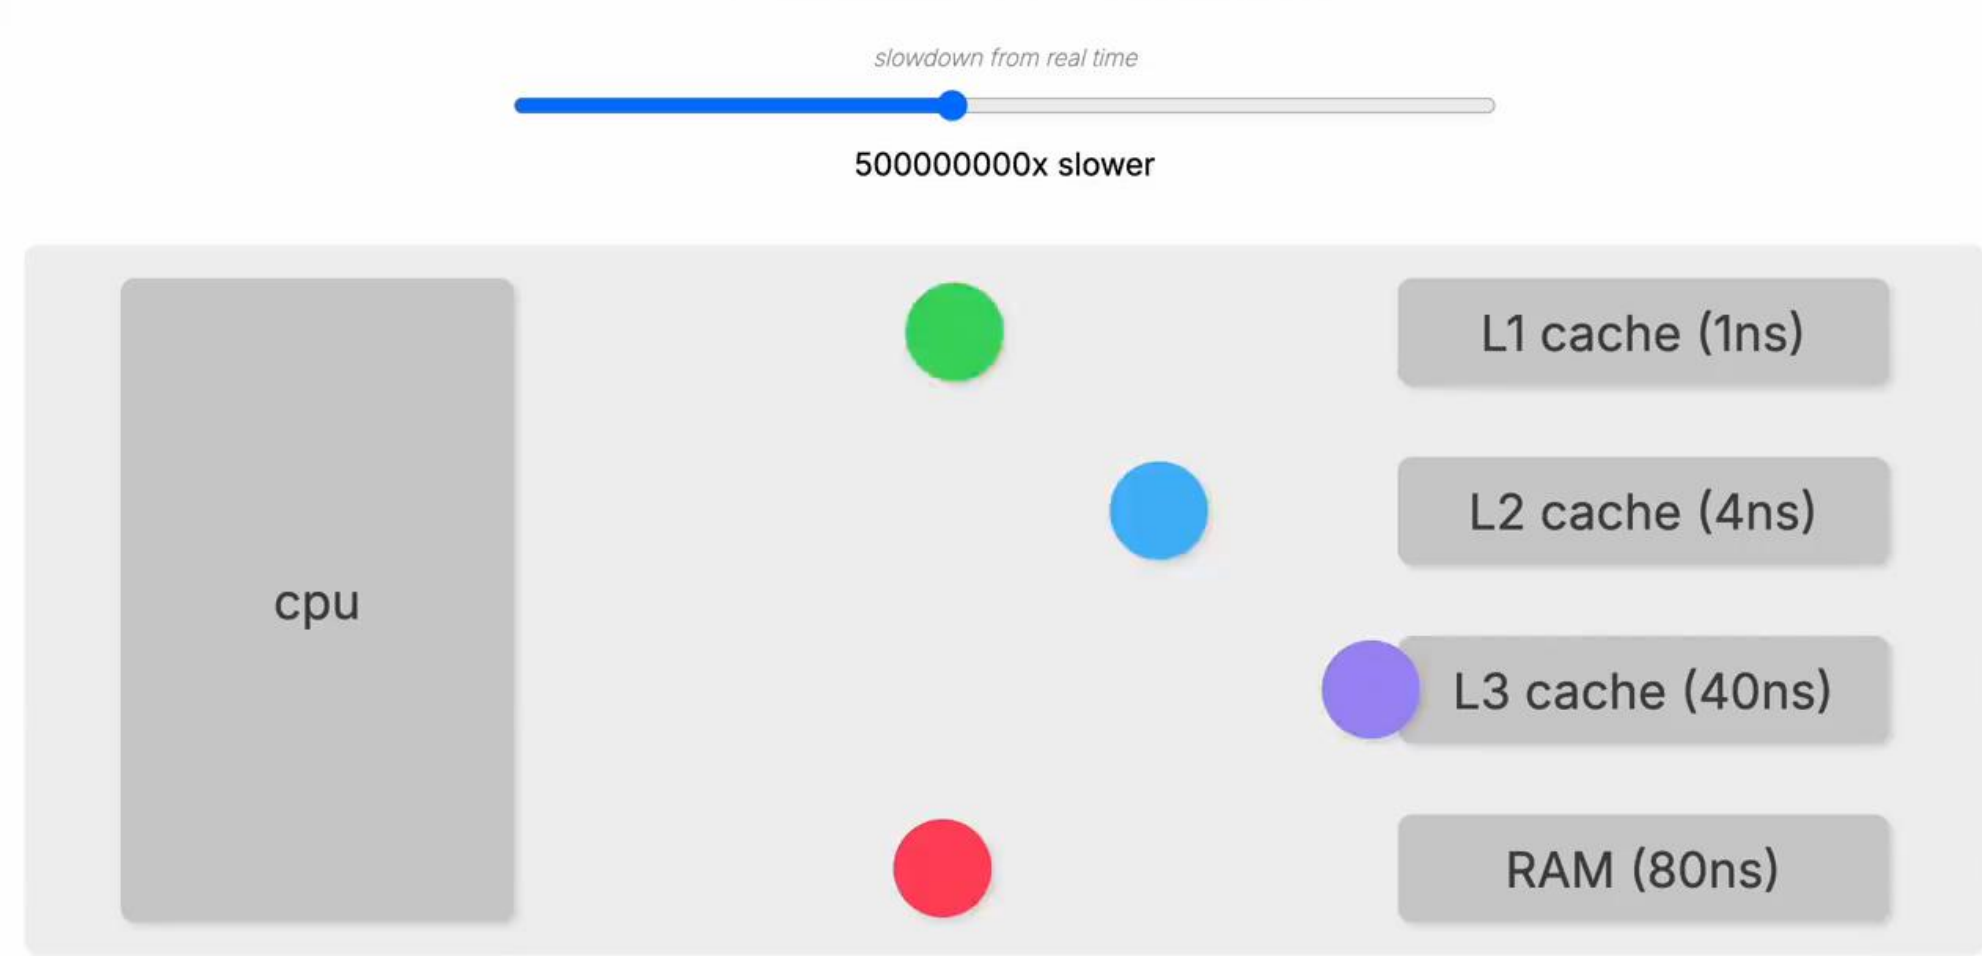


[Hennessy & Patterson, 2017]



内存层次

CPU Cache and RAM





## 延迟

**Table 2.2** Example Time Scale of System Latencies

Event	Latency	Scaled
1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access (DRAM, from CPU)	120 ns	6 min
Solid-state disk I/O (flash memory)	50–150 $\mu$ s	2–6 days
Rotational disk I/O	1–10 ms	1–12 months
Internet: San Francisco to New York	40 ms	4 years
Internet: San Francisco to United Kingdom	81 ms	8 years
Internet: San Francisco to Australia	183 ms	19 years
TCP packet retransmit	1–3 s	105–317 years
OS virtualization system reboot	4 s	423 years
SCSI command time-out	30 s	3 millennia
Hardware (HW) virtualization system reboot	40 s	4 millennia
Physical system reboot	5 m	32 millennia

## What Compilers Cannot Do



**Sebastian Aaltonen**

@SebAaltonen

...

It's also important to emphasize that the compiler is not allowed to change your data layout. Modern CPUs running at 5 GHz with 6-wide ALU pipes and GPUs with dual float pipes can handle couple of extra ALUs. Memory accesses are the bottleneck, and compiler can't help with it.

## Guidance

For data:

- **Where practical, employ linear array traversals.**
  - ➔ “I don’t know [data structure], but I know an array will beat it.”
- **Use as much of a cache line as possible.**
  - ➔ Bruce Dawson’s antipattern (from reviews of video games):

```
struct Object {                // assume sizeof(Object) ≥ 64
    bool isLive;                // possibly a bit field
    ...
};

std::vector<Object> objects;    // or an array

for (std::size_t i = 0; i < objects.size(); ++i) { // pathological if
    if (objects[i].isLive)    // most objects
        doSomething();        // not alive
}
```
- **Be alert for false sharing in MT systems.**

## Example: Loop interchange

### 矩阵乘法

*// Column-major order*


```
for (row = 0; row < NROWS; row++)  
  for (col = 0; col < NCOLS; col++)  
    matrix[col][row] = row + col;
```

=>


*// Row-major order*

```
for (row = 0; row < NROWS; row++)  
  for (col = 0; col < NCOLS; col++)  
    matrix[row][col] = row + col;
```

column-major



row-major



## Example: Loop interchange

矩阵乘法

```
// Multiply two square matrices
void multiply(Matrix &result, const Matrix &a,
             const Matrix &b) {
    zero(result);
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            for (int k = 0; k < N; k++) {
                result[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```

# Loop interchange

## 矩阵乘法

```
// Multiply two square matrices
void multiply(Matrix &result, const Matrix &a,
const Matrix &b) {
    zero(result);
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            for (int k = 0; k < N; k++) {
                result[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```

Elapsed Time<sup>?</sup>: **0.604s** 🚩

⌕ Clockticks:	2,768,832,000	
⌕ Instructions Retired:	2,773,056,000	↗
⌕ CPI Rate <sup>?</sup> :	0.998	
⌕ MUX Reliability <sup>?</sup> :	<b>0.676</b> 🚩	
⌕ P-Core:		
⌕ Retiring <sup>?</sup> :	28.6%	of Pipeline Slots
⌕ Front-End Bound <sup>?</sup> :	14.3%	of Pipeline Slots
⌕ Bad Speculation <sup>?</sup> :	0.0%	of Pipeline Slots
⌕ Back-End Bound <sup>?</sup> :	69.6%	of Pipeline Slots

↘

# Loop interchange

## 矩阵乘法

```
// Multiply two square matrices
void multiply(Matrix &result, const Matrix &a,
const Matrix &b) {
    zero(result);
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            for (int k = 0; k < N; k++) {
                result[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```


顺序访问      顺序访问      跳N访问

Back-End Bound	69.6%	of Pipeline Slots
Memory Bound	23.2%	of Pipeline Slots
L1 Bound	24.3%	of Clockticks
DTLB Overhead	97.1%	of Clockticks
Loads Blocked by Store Forwarding	0.0%	of Clockticks
Lock Latency	0.0%	of Clockticks
Split Loads	0.0%	of Clockticks
4K Aliasing	0.0%	of Clockticks
FB Full	0.0%	of Clockticks
L2 Bound	0.0%	of Clockticks
L3 Bound	0.9%	of Clockticks
DRAM Bound	0.0%	of Clockticks
Store Bound	0.0%	of Clockticks
Core Bound	46.4%	of Pipeline Slots

# Loop interchange

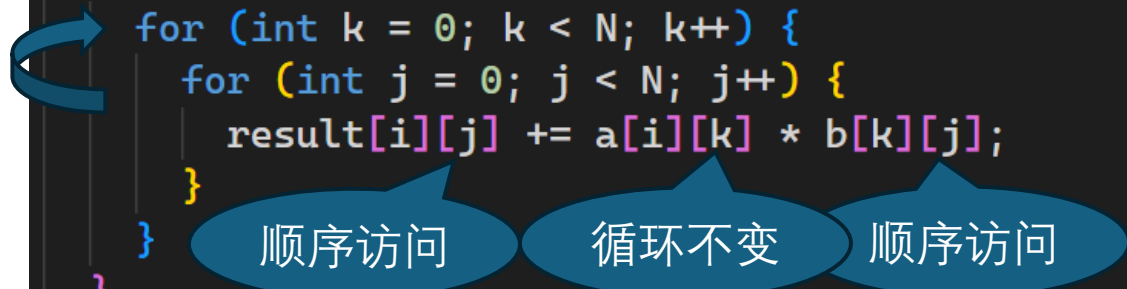
## 优化前

```
// Multiply two square matrices
void multiply(Matrix &result, const Matrix &a,
const Matrix &b) {
    zero(result);
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            for (int k = 0; k < N; k++) {
                result[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```



## 优化后

```
// Multiply two square matrices
void multiply(Matrix &result, const Matrix &a,
const Matrix &b) {
    zero(result);
    for (int i = 0; i < N; i++) {
        for (int k = 0; k < N; k++) {
            for (int j = 0; j < N; j++) {
                result[i][j] += a[i][k] * b[k][j];
            }
        }
    }
}
```



**88% speed up**

[https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/memory\\_bound/loop\\_interchange\\_1](https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/memory_bound/loop_interchange_1)



# Loop interchange

Source	Back-End Bound					
	Memory Bound					Core Bound
	L1 Bound	L2 Bound	L3 Bound	DRAM Bound	Store Bound	
for (int i = 0; i < N; i++) {						
for (int j = 0; j < N; j++) {						
for (int k = 0; k < N; k++) {	0.0%	0.0%	0.0%	0.0%	0.0%	
result[i][j] += a[i][k] * b[k][j];	28.4%	0.0%	0.0%	0.0%	0.0%	100.0%
}						
}						
}						

Source	Back-End Bound					
	Memory Bound					Core Bound
	L1 Bound	L2 Bound	L3 Bound	DRAM Bound	Store Bound	
for (int i = 0; i < N; i++) {						
for (int k = 0; k < N; k++) {	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%
for (int j = 0; j < N; j++) {	0.0%	0.0%	0.0%	0.0%	0.0%	0.0%
result[i][j] += a[i][k] * b[k][j];	0.0%	0.0%	0.0%	0.0%	0.0%	
}						
}						
}						

# Support 16 KB page sizes

## Benefits and performance gains

Devices configured with 16 KB page sizes use slightly more memory on average, but also gain various performance improvements for both the system and apps:

- Lower app launch times while the system is under memory pressure: 3.16% lower on average, with more significant improvements (up to 30%) for some apps that we tested
- Reduced power draw during app launch: 4.56% reduction on average
- Faster camera launch: 4.48% faster hot starts on average, and 6.60% faster cold starts on average
- Improved system boot time: improved by 8% (approximately 950 milliseconds) on average

These improvements are based on our initial testing, and results on actual devices will likely differ. We'll provide additional analysis of potential gains for apps as we continue our testing.

## Hardware Effects

### Memory Order Violation

```
std::array<uint32_t, 256> hist;  
hist.fill(0);  
for (int i = 0; i < image.width * image.height; ++i)  
|   hist[image.data[i]]++;  
return hist;
```

求灰度图的直方图

0xFF 0xFF 0xFF 0xFF 0xFF 0xFF ...

Then all updates to hist[0xFF] will be serialized.

# Hardware Effects

## Memory Order Violation

```
for (int i = 0; i < image.width * image.height; ++i)  
    hist[image.data[i]]++;
```

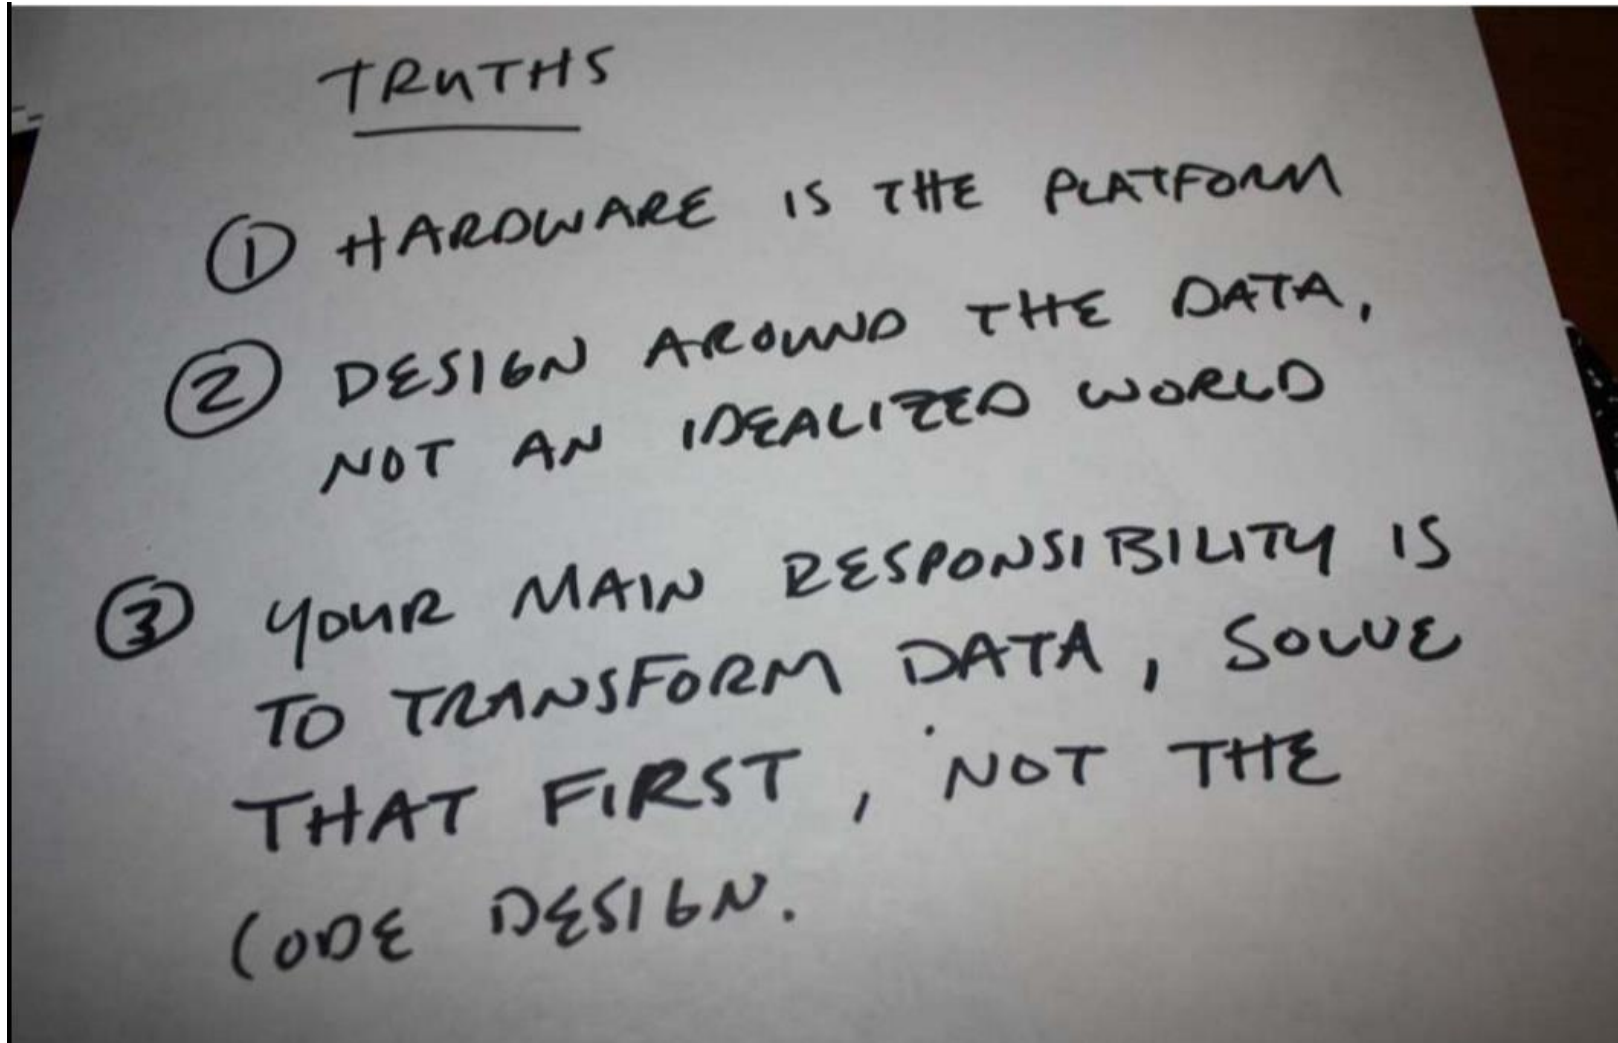


```
for (; i + 3 < image.width * image.height; i += 4) {  
    hist1[image.data[i+0]]++;  
    hist2[image.data[i+1]]++;  
    hist3[image.data[i+2]]++;  
    hist4[image.data[i+3]]++;  
}
```

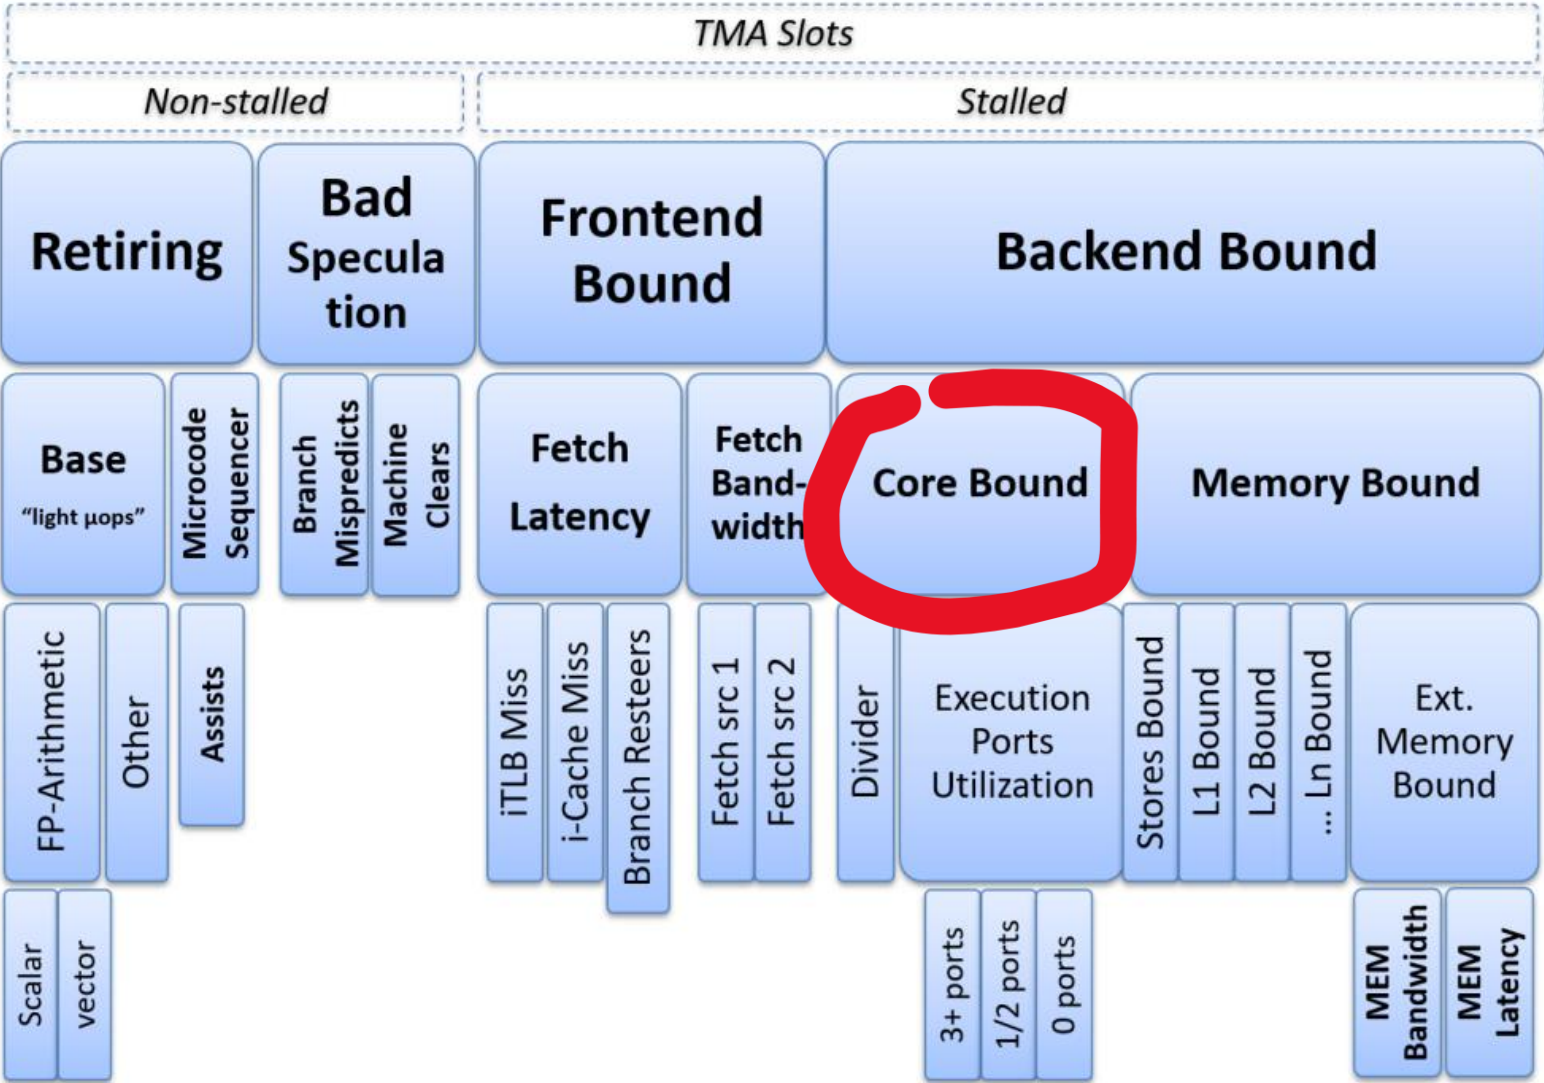
**0-60% speed up**

[https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/memory\\_bound/mem\\_order\\_violation\\_1](https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/memory_bound/mem_order_violation_1)

## Data Oriented Design




# Four Corners



# Data Dependencies

## Linked List



```
while (n) {  
    sum += n→val;  
    n = n→next;  
}
```

Pointer chasing does not benefit from OOO execution and thus will run at the speed of an in-order CPU.

# Data Dependencies

## Simulate random particle movement

```
struct Particle {
    float x; float y; float velocity;
};

class XorShift32 {
    uint32_t val;
public:
    XorShift32 (uint32_t seed) : val(seed) {}
    uint32_t gen() {
        val ^= (val << 13);
        val ^= (val >> 17);
        val ^= (val << 5);
        return val;
    }
};

static float sine(float x) {
    const float B = 4 / PI_F;
    const float C = -4 / (PI_F * PI_F);
    return B * x + C * x * std::abs(x);
}

static float cosine(float x) {
    return sine(x + (PI_F / 2));
}
```

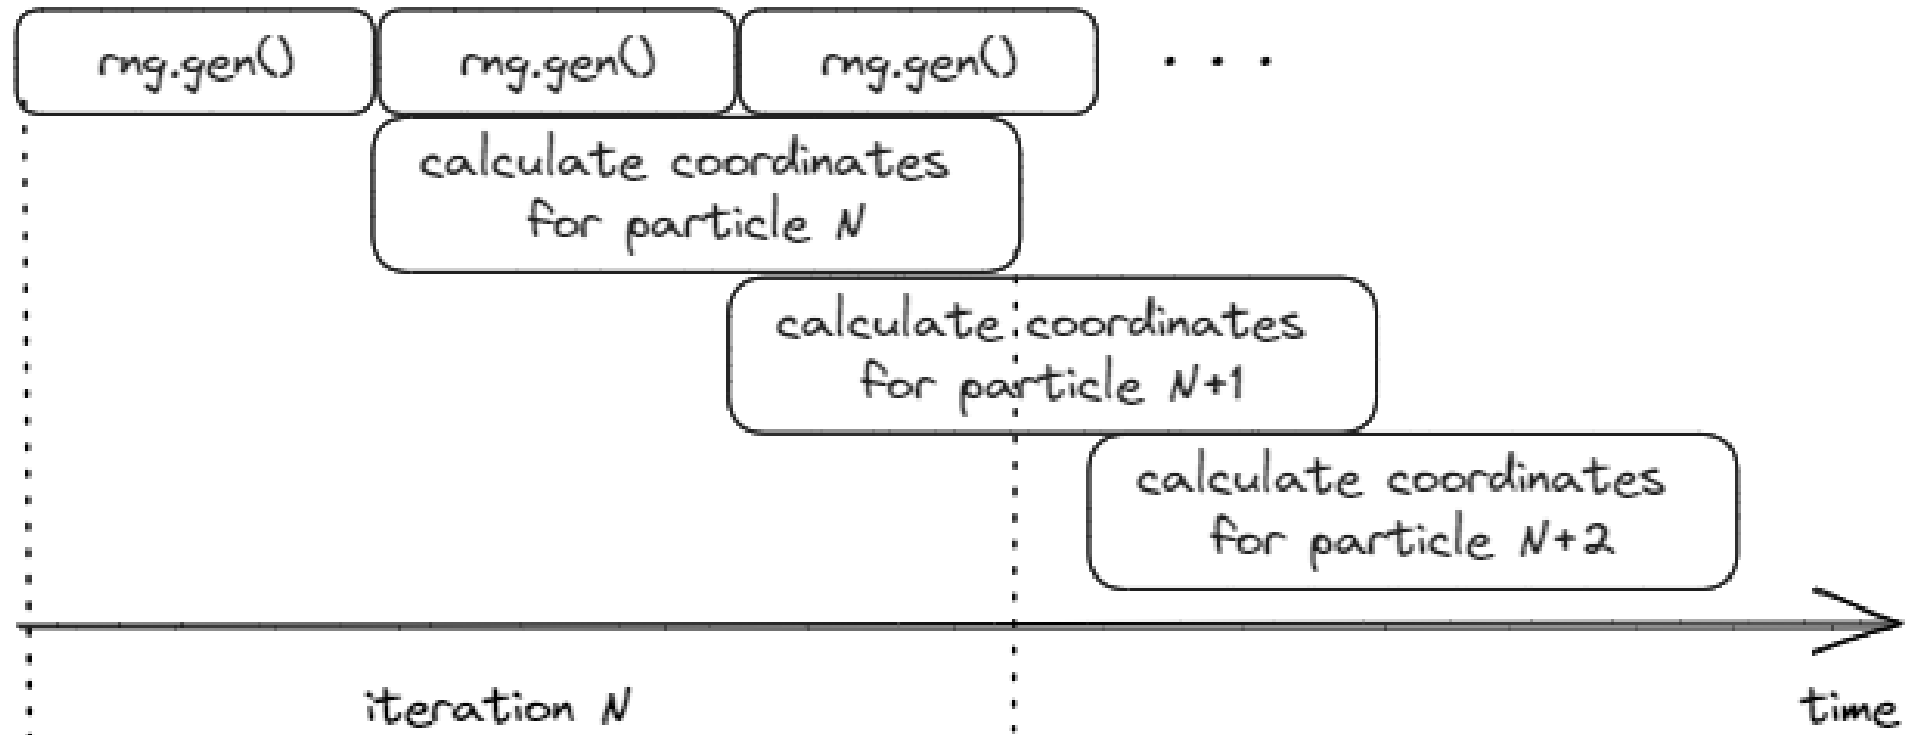
```
/* Map degrees [0;UINT32_MAX) to radians [0;2*pi)*/
float DEGREE_TO_RADIAN = (2 * PI_D) / UINT32_MAX;

void particleMotion(vector<Particle> &particles,
                    uint32_t seed) {
    XorShift32 rng(seed);
    for (int i = 0; i < STEPS; i++)
        for (auto &p : particles) {
            uint32_t angle = rng.gen();
            float angle_rad = angle * DEGREE_TO_RADIAN;
            p.x += cosine(angle_rad) * p.velocity;
            p.y += sine(angle_rad) * p.velocity;
        }
}
```



# Data Dependencies

## Visualization of dependent execution



# Data Dependencies

## Solution

```
void particleMotion(vector<Particle> &particles,
                   uint32_t seed1, uint32_t seed2) {
    XorShift32 rng1(seed1);
    XorShift32 rng2(seed2);
    for (int i = 0; i < STEPS; i++) {
        for (int j = 0; j + 1 < particles.size(); j += 2) {
            uint32_t angle1 = rng1.gen();
            float angle_rad1 = angle1 * DEGREE_TO_RADIAN;
            particles[j].x += cosine(angle_rad1) * particles[j].velocity;
            particles[j].y += sine(angle_rad1) * particles[j].velocity;
            uint32_t angle2 = rng2.gen();
            float angle_rad2 = angle2 * DEGREE_TO_RADIAN;
            particles[j+1].x += cosine(angle_rad2) * particles[j+1].velocity;
            particles[j+1].y += sine(angle_rad2) * particles[j+1].velocity;
        }
        // remainder (not shown)
    }
}
```

Once you do this transformation, the compiler starts autovectorizing the body of the loop.

**55% speed up**

## Inline



# Function call overhead

```
void foo(some_args) {  
    //...  
    bar();  
    // CALL bar  
    //...  
}
```

```
void bar() {  
    // PUSH rbp # prologue  
    // PUSH rbx  
    //  
    // ... # function body  
    //  
    // POP rbx # epilogue  
    // POP rbp  
    // RET  
}
```

# Inline

```
static int compare(const void *lhs, const void *rhs) {
    auto &a = *reinterpret_cast<const S *>(lhs);
    auto &b = *reinterpret_cast<const S *>(rhs);

    if (a.key1 < b.key1)
        return -1;

    if (a.key1 > b.key1)
        return 1;

    if (a.key2 < b.key2)
        return -1;

    if (a.key2 > b.key2)
        return 1;

    return 0;
}

void solution(std::array<S, N> &arr) {
    qsort(arr.data(), arr.size(), sizeof(S), compare);
}
```



```
void solution(std::array<S, N> &arr) {
    std::sort(arr.begin(), arr.end(), [](S& a, S& b)
    {
        return a.key1 < b.key1 ||
            (a.key1 == b.key1) && (a.key2 < b.key2);
    });
}
```

**47% speed up**

[https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/core\\_bound/function\\_inlining\\_1](https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/core_bound/function_inlining_1)

## Auto vectorization: Rely on Compiler

Does the following loop vectorize?

```
void daxpy4(double *__restrict z, double a,  
            const double *__restrict x,  
            const double *__restrict y,  
            size_t n) {  
    for (size_t i = 0; i < n; i += 4) {  
        z[i]    = a * x[i]    + y[i];  
        z[i+1]  = a * x[i+1] + y[i+1];  
        z[i+2]  = a * x[i+2] + y[i+2];  
        z[i+3]  = a * x[i+3] + y[i+3];  
    }  
    // ...  
}
```


# What Compilers Cannot Do

## 自动向量化失败

```
remark: loop not vectorized: could not determine number of  
loop iterations [-Rpass-analysis=loop-vectorize] x86-64 clang  
12.0.0 #2
```

```
remark: loop not vectorized [-Rpass-missed=loop-  
vectorize] x86-64 clang 12.0.0 #2
```

```
No quick fixes available
```



```
for (size_t i = 0; i < n; i += 4) {  
    z[i]    = a * x[i]    + y[i];  
    z[i+1]  = a * x[i+1]  + y[i+1];  
    z[i+2]  = a * x[i+2]  + y[i+2];  
    z[i+3]  = a * x[i+3]  + y[i+3];  
}  
// ...  
}
```

clang12 -O3 -Rpass-analysis=loop-vectorize -Rpass=loop-vectorize -Rpass-missed=loop-vectorize

# What Compilers Cannot Do

## 编译器的限制(bug)

```
void daxpy4(double * restrict z, double a,  
            const unsigned __restrict x,  
            const int64_t __restrict y,  
            size_t n) {  
    for (size_t i = 0; i < n; i += 4) {  
        z[i]    = a * x[i]    + y[i];  
        z[i+1]  = a * x[i+1] + y[i+1];  
        z[i+2]  = a * x[i+2] + y[i+2];  
        z[i+3]  = a * x[i+3] + y[i+3];  
    }  
    // ...  
}
```

**PROBLEM:** In C, the behavior of unsigned-integer overflow is defined to wrap around.

# What Compilers Cannot Do

## 自动向量化

```
#inc remark: the cost-model indicates that interleaving is not
#inc beneficial [-Rpass-analysis=loop-vectorize] x86-64 clang
12.0.0 #2

void remark: vectorized loop (vectorization width: 2, interleaved
count: 1) [-Rpass=loop-vectorize] x86-64 clang 12.0.0 #2

No quick fixes available

for (int64_t i = 0; i < n; i += 4) {
    z[i] = a * x[i] + y[i];
    z[i+1] = a * x[i+1] + y[i+1];
    z[i+2] = a * x[i+2] + y[i+2];
    z[i+3] = a * x[i+3] + y[i+3];
}
// ...
}
```

int64\_t

**SOLUTION:** Use signed integer types, signed-integer overflow has undefined behavior.

Or upgrade clang



## What Compilers Cannot Do

float

```
float calcSum(float* a, unsigned N) {  
    float sum = 0.0f;  
    for (unsigned i = 0; i < N; i++) {  
        sum += a[i];  
    }  
    return sum;  
}
```

SOLUTION:

-ffast-math

remark: loop not vectorized: cannot prove it is safe to reorder floating-point operations; allow reordering by specifying '#pragma clang loop vectorize(enable)' before the loop or by providing the compiler option '-ffast-math' [-Rpass-analysis=loop-vectorize] x86-64 clang 19.1.0 (assertions) #2

# Compiler Intrinsics

## When the compiler fails

```
// a.cpp
float calcSum(float* a, unsigned N) {
    float sum = 0.0f;
    for (unsigned i = 0; i < N; i++) {
        sum += a[i];
    }
    return sum;
}
```



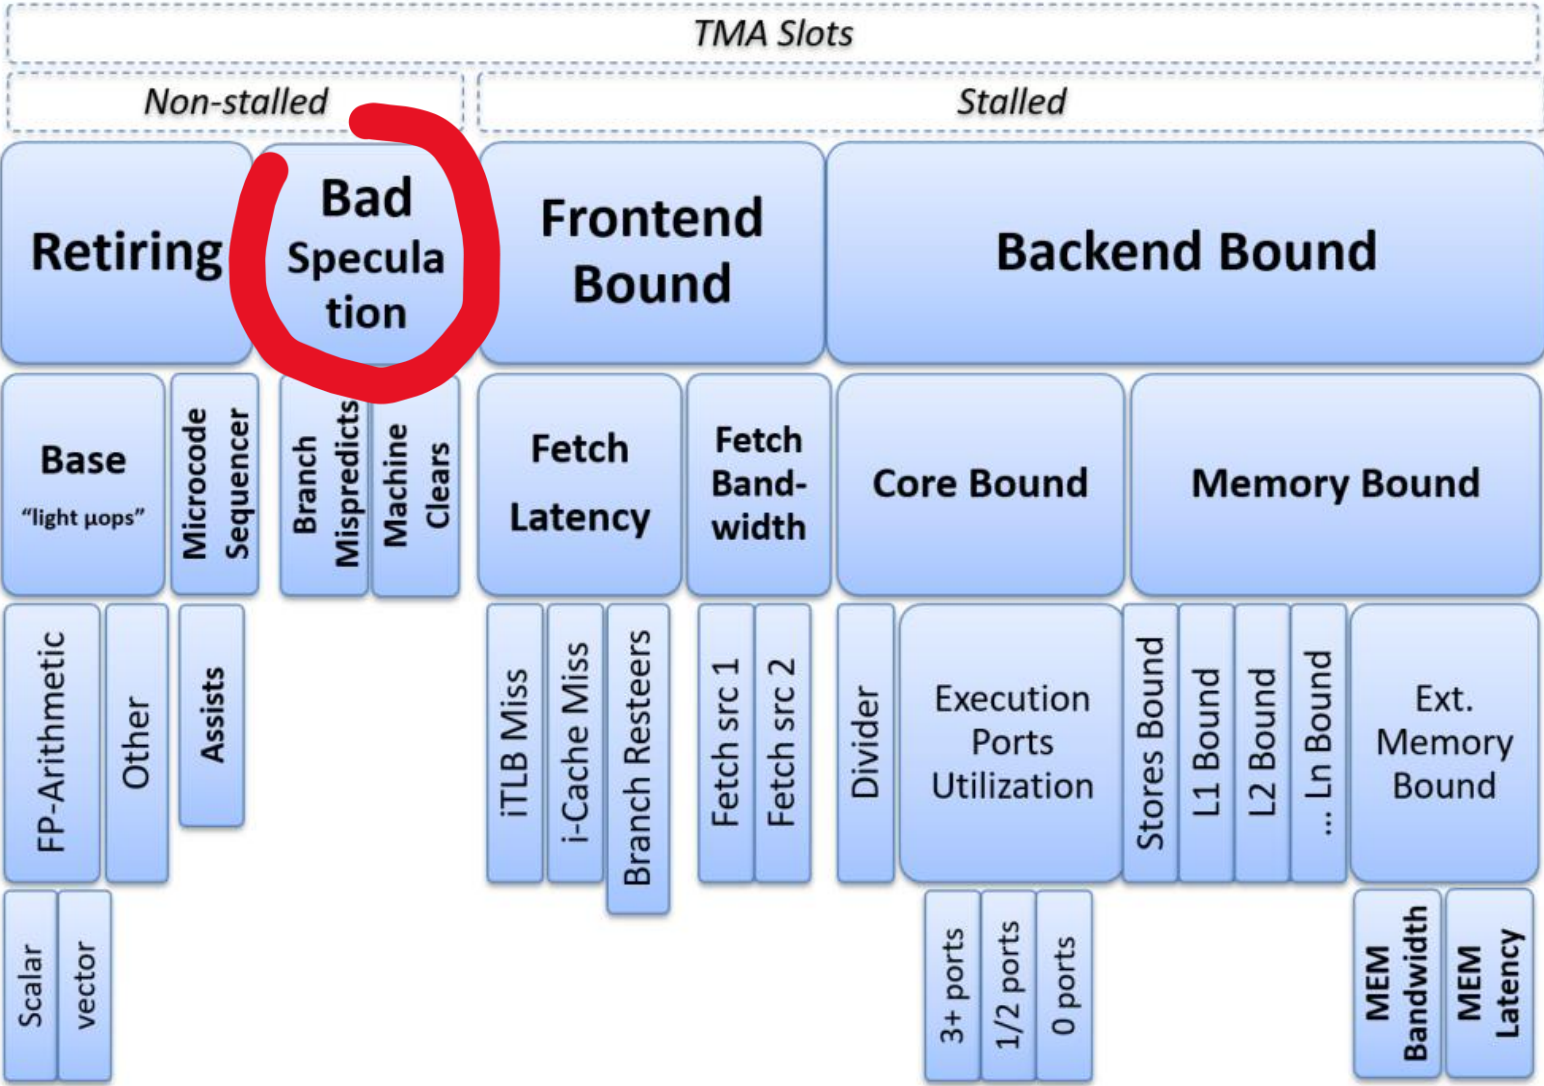
```
#include <immintrin.h>

float calcSum(float* a, unsigned N) {
    __m128 sum = _mm_setzero_ps();      // init sum with zeros
    unsigned i = 0;
    for (; i + 3 < N; i += 4) {
        __m128 vec = _mm_loadu_ps(a + i); // load 4 floats from array
        sum = _mm_add_ps(sum, vec);      // accumulate vec into sum
    }

    // Horizontal sum of the 128-bit vector
    __m128 shuf = _mm_movehdup_ps(sum); // broadcast elements 3,1 to 2,0
    sum = _mm_add_ps(sum, shuf);        // partial sums [0+1] and [2+3]
    shuf = _mm_movehl_ps(shuf, sum);    // high half -> low half
    sum = _mm_add_ss(sum, shuf);        // result in the lower element
    float result = _mm_cvtss_f32(sum);  // nop (compiler eliminates it)

    // Process any remaining elements
    for (; i < N; i++)
        result += a[i];
    return result;
}
```

# Four Corners



# 生命游戏

规则：

1. 每个细胞有两种状态 - 存活或死亡，每个细胞与以自身为中心的周围八格细胞产生互动
2. 当前细胞为存活状态时，当周围的存活细胞低于2个时（不包含2个），该细胞变成死亡状态。
3. 当前细胞为存活状态时，当周围有2个或3个存活细胞时，该细胞保持原样。
4. 当前细胞为存活状态时，当周围有超过3个存活细胞时，该细胞变成死亡状态。
5. 当前细胞为死亡状态时，当周围有3个存活细胞时，该细胞变成存活状态。



# 生命游戏

```
switch(aliveNeighbours) {  
    // 1. Cell is lonely and dies  
    case 0:  
    case 1:  
        future[i][j] = 0;  
        break;  
    // 2. Remains the same  
    case 2:  
        future[i][j] = current[i][j];  
        break;  
    // 3. A new cell is born  
    case 3:  
        future[i][j] = 1;  
        break;  
    // 4. Cell dies due to over population  
    default:  
        future[i][j] = 0;  
}
```

⌵ P-Core:	
⌵ Retiring ?:	18.2% of Pipeline Slots
⌵ Front-End Bound ?:	13.8% of Pipeline Slots
⌵ Bad Speculation ?:	58.9% 🚩 of Pipeline Slots
Branch Mispredict ?:	51.6% 🚩 of Pipeline Slots
Machine Clears ?:	7.3% of Pipeline Slots
⌵ Back-End Bound ?:	9.1% of Pipeline Slots

## Branch => cmove

```
switch(aliveNeighbours) {  
    // 1. Cell is lonely and dies  
    case 0:  
    case 1:  
        future[i][j] = 0;  
        break;  
    // 2. Remains the same  
    case 2:  
        future[i][j] = current[i][j];  
        break;  
    // 3. A new cell is born  
    case 3:  
        future[i][j] = 1;  
        break;  
    // 4. Cell dies due to over population  
    default:  
        future[i][j] = 0;  
}
```



```
int cell = current[i][j];  
if (__builtin_unpredictable(aliveNeighbours != 2))  
    cell = 0;  
if (__builtin_unpredictable(aliveNeighbours == 3))  
    cell = 1;  
future[i][j] = cell;
```

**49% speed up**

## 另一个热点

### 边界判断

```
// finding the number of neighbours that are alive
for(int p = -1; p ≤ 1; p++) // row-offset (-1,0,1)
    for(int q = -1; q ≤ 1; q++) // col-offset (-1,0,1)
        if((i + p < 0) || // if row offset less than UPPER boundary
            (i + p > M - 1) || // if row offset more than LOWER boundary
            (j + q < 0) || // if column offset less than LEFT boundary
            (j + q > N - 1)) // if column offset more than RIGHT boundary
            continue;
        aliveNeighbours += current[i + p][j + q];
    }
}
```

消除这4个判断

Solution: 上下左右的边界各往外扩一格即可

**93% speed up**

[https://github.com/jsjxietian/perf-ninja-solution/tree/master/labs/bad\\_speculation/branches\\_to\\_cmov\\_1](https://github.com/jsjxietian/perf-ninja-solution/tree/master/labs/bad_speculation/branches_to_cmov_1)

# Lookup table

```
static std::size_t mapToBucket(std::size_t v) {  
    // size of a bucket  
    if (v < 13) return 0; // 13  
    else if (v < 29) return 1; // 16  
    else if (v < 41) return 2; // 12  
    else if (v < 53) return 3; // 12  
    else if (v < 71) return 4; // 18  
    else if (v < 83) return 5; // 12  
    else if (v < 100) return 6; // 17  
    return DEFAULT_BUCKET;  
}
```



```
const static int buckets[101] = {  
    0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, // thirteen 0s  
    1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, // sixteen 1s  
    2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, // twelve 2s  
    3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, // twelve 3s  
    4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, // eighteen 4s  
    5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, // twelve 5s  
    6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, // seventeen 6s  
    DEFAULT_BUCKET  
};  
  
static std::size_t mapToBucket(std::size_t v) {  
    constexpr auto Nelements = sizeof (buckets) / sizeof (int);  
    return buckets[std::min(v, Nelements - 1)];  
}
```

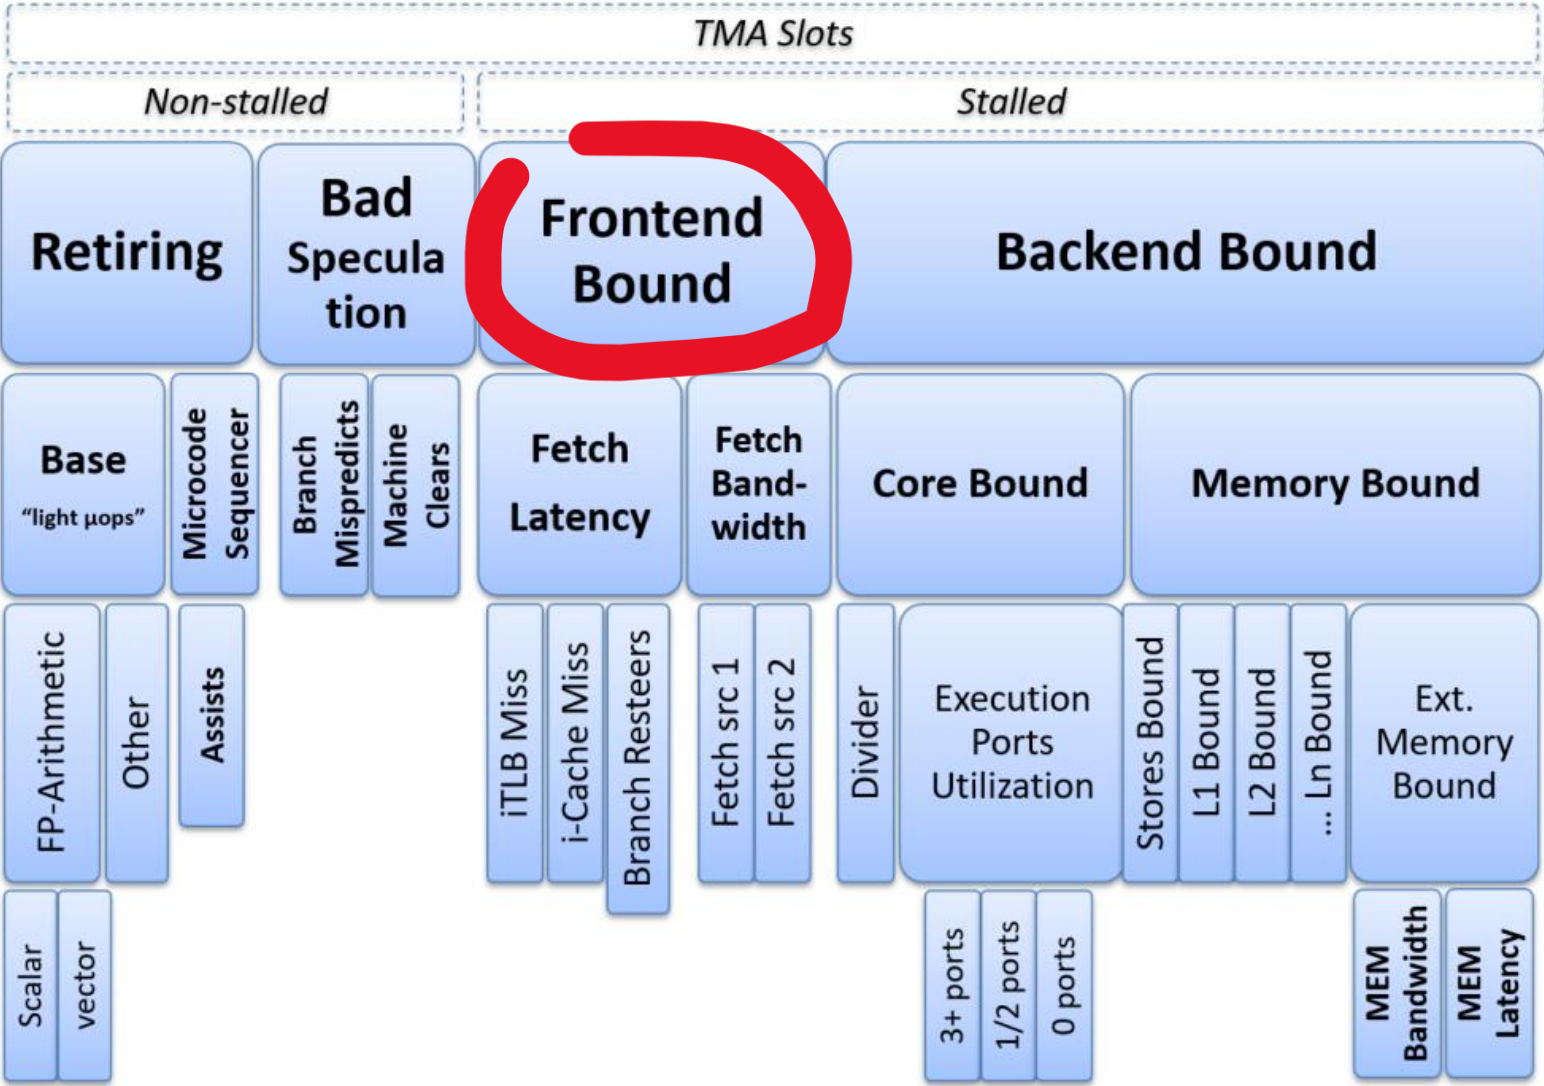
**88% speed up**

[https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/bad\\_speculation/lookup\\_tables\\_1](https://github.com/jsjtxietian/perf-ninja-solution/tree/master/labs/bad_speculation/lookup_tables_1)


Note: always measure, 右边的增加了cache的压力



# Four Corners



# Machine Code Layout Causes Noise



## Machine Code Layout Causes Noise

# Layout biases measurement

Mytkowicz et al. (ASPLOS'09)

## Link Order

Changes function addresses

**Larger than the  
impact of -O3**

## Environment

## Variable Size

Moves the program stack

# Machine Code Layout

## 抖音研发实践：基于二进制文件重排的解决方案 APP启动速度提升超15%

原创 Leo 字节跳动技术团队 2019年08月09日 19:43

### 背景

---

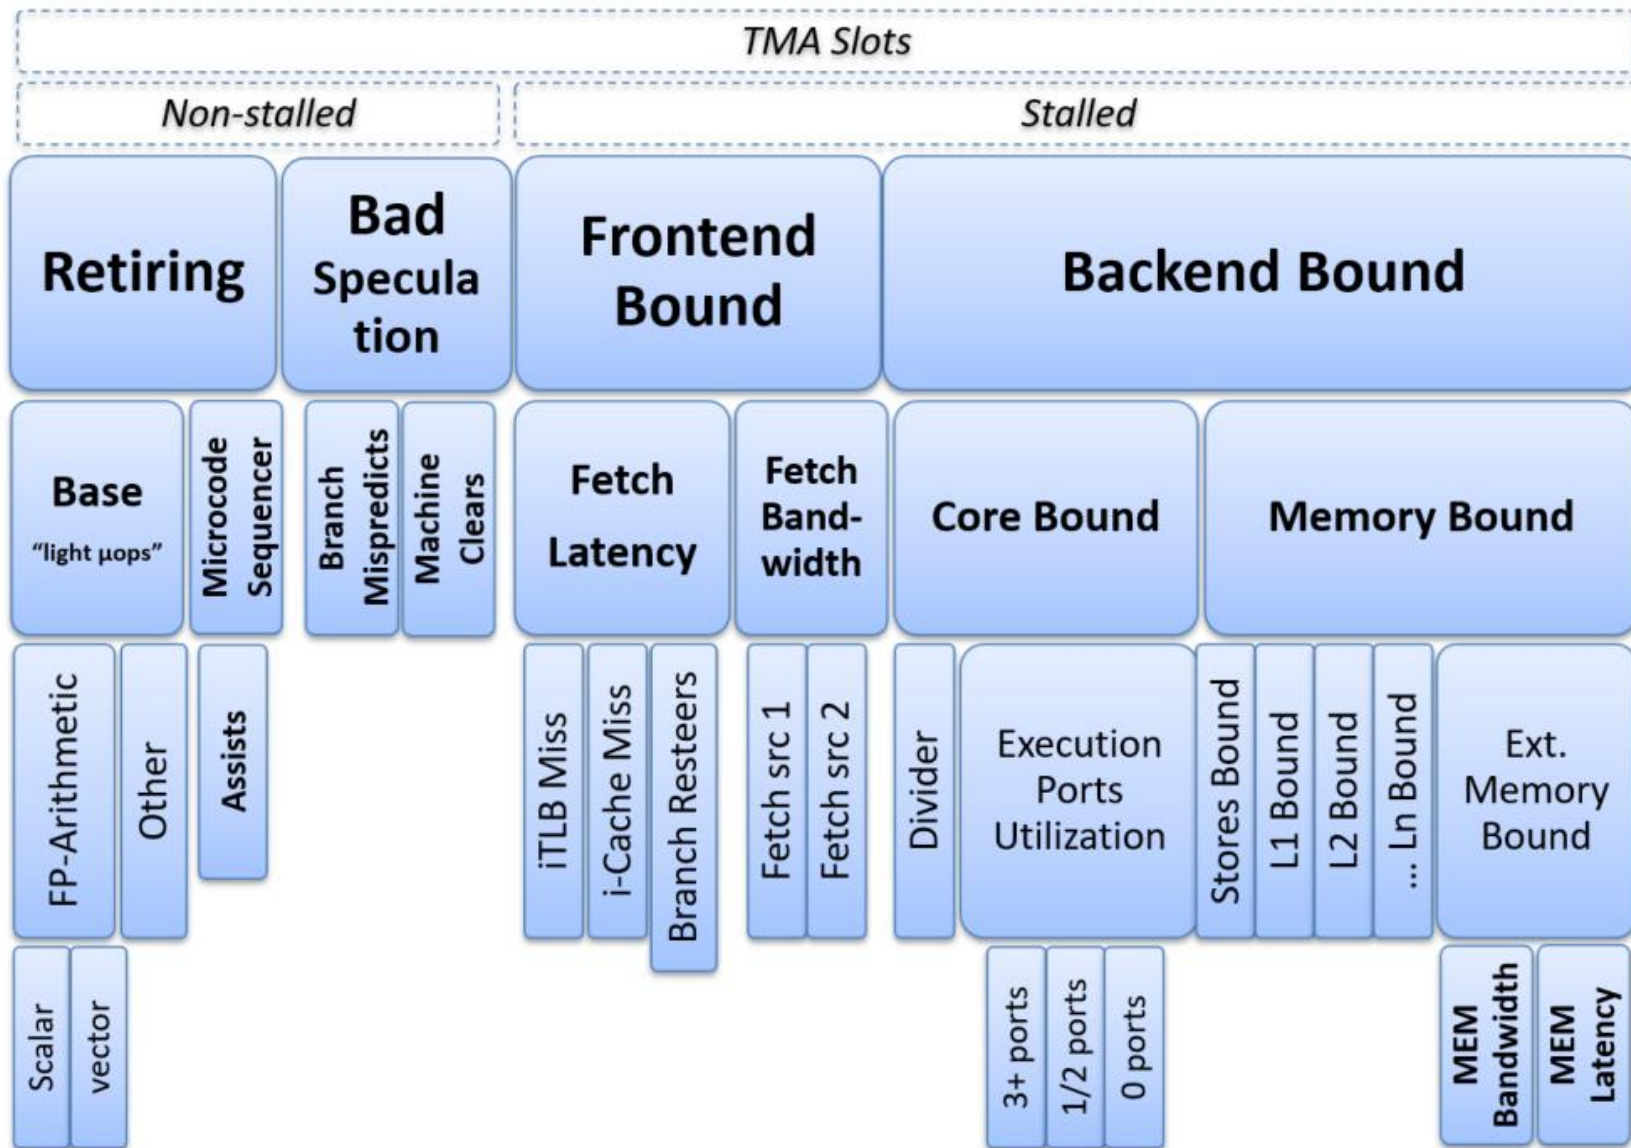
启动是App给用户的第一印象，对用户体验至关重要。抖音的业务迭代迅速，如果放任不管，启动速度会一点点劣化。为此抖音iOS客户端团队做了大量优化工作，除了传统的修改业务代码方式，我们还做了些开拓性的探索，发现修改代码在二进制文件的布局可以提高启动性能，方案落地后在抖音上启动速度提高了约15%。

本文从原理出发，介绍了我们是如何通过静态扫描和运行时trace找到启动时候调用的函数，然后修改编译参数完成二进制文件的重新排布。

# Machine Code Layout

优化一个Page Fault，启动速度提升0.6~0.8ms






**Final**

# Always Measure!

1. Mental model can never be as accurate as the actual microarchitecture design of a CPU. Always measure!
2. When measuring performance, understand the underlying technical reasons for the performance results .

More

Benchmark



Perf

